

Certainly! Below are some **Linux interview questions** tailored for a DevOps engineer with **3 years of experience**. These questions cover a variety of topics, including basic Linux commands, system administration, file management, networking, and performance tuning.

1. What are the key differences between Linux and Windows operating systems?

Answer:

- **Kernel:** Linux uses a monolithic kernel, whereas Windows uses a hybrid kernel.
- **File System:** Linux uses file systems like ext4, XFS, while Windows uses NTFS.
- **Permissions:** Linux uses **user** and **group** permissions for files (read, write, execute), whereas Windows uses **ACLs**.
- **Command-line Interface (CLI):** Linux heavily relies on CLI (Bash), while Windows uses PowerShell or Command Prompt.
- **Package Management:** Linux uses package managers like **apt**, **yum**, and **dnf**, while Windows uses MSI or Windows Store.

2. What are the different types of file permissions in Linux?

Answer:

Linux file permissions can be read, write, and execute: - **r** (read): Grants permission to view the contents of the file. - **w** (write): Grants permission to modify the file. - **x** (execute): Grants permission to run the file as a program or script.

Permissions are assigned to three entities: - **Owner:** The user who owns the file. - **Group:** The group that the file belongs to. - **Others:** Everyone else.

Example: `bash chmod 755 myfile.txt`

This means the owner can read/write/execute, while the group and others can read/execute.

3. What is the purpose of the chmod command?

Answer:

The `chmod` command is used to change the permissions of a file or directory. It can modify the read, write, and execute permissions for the owner, group, and others. Permissions can be changed using symbolic or numeric modes.

Example (symbolic mode): `bash chmod u+x myscript.sh # Adds execute permission to the user`

Example (numeric mode): `bash chmod 644 myfile.txt # Owner can read/write, others can only read`

4. Explain the difference between soft links and hard links in Linux.

Answer:

- **Hard Link:** Points directly to the data on disk. It creates a second directory entry for the same file, sharing the same inode number. Deleting one hard link does not remove the data as long as another hard link exists. - **Soft Link (Symbolic Link):** Points to the path of a file or directory. It has its own inode and references the original file by name. If the original file is deleted, the symbolic link becomes broken.

Example of creating a hard link: `bash ln original_file hard_link`

Example of creating a soft link: `bash ln -s original_file soft_link`

5. What are the basic differences between a process and a thread in Linux?

Answer: - **Process:** A process is an independent program that runs in its own memory space. Processes are isolated from each other. - **Thread:** A thread is a smaller unit of execution within a process. Threads within the same process share the same memory space, which allows them to communicate more easily but also makes them more susceptible to errors such as race conditions.

6. What is the difference between kill and killall in Linux?

Answer: - **kill:** Sends a signal (default is SIGTERM) to a specific process by its process ID (PID). Example:

`bash kill 1234 # Terminates the process with PID 1234` -

killall: Sends a signal to all processes with a given name. Example:

`bash killall apache2 # Terminates all processes with the name 'apache2'`

7. How do you check system resource usage (CPU, memory, disk space) on a Linux machine?

Answer: - **CPU Usage:**

`bash top # Displays real-time CPU usage mpstat`
`# Provides CPU statistics`

- **Memory Usage:**

`bash free -m # Shows memory usage in MB vmstat`
`# Provides virtual memory statistics`

- **Disk Usage:**

`bash df -h # Displays disk space usage for all mounted filesystems`
`du -sh # Displays the size of a specific directory`

8. What is the `ps` command in Linux?

Answer:

The `ps` (process status) command shows information about the currently running processes on the system. You can use it with different options to get detailed process information.

Example: `bash ps aux # Lists all running processes` `ps -ef # Shows a full-format list of processes`

9. How do you configure a static IP address on a Linux system?

Answer: To configure a static IP address on a Linux system, you can modify the network configuration file for your interface. For example, on Ubuntu, you can modify the `/etc/netplan/*.yaml` file.

Example: `yaml network: version: 2 renderer: networkd`
`ethernets: eth0: dhcp4: false addresses:`
`- 192.168.1.100/24 gateway4: 192.168.1.1 nameservers:`
`addresses: - 8.8.8.8 - 8.8.4.4`

Then apply the configuration: `bash sudo netplan apply`

10. What is `cron` and how do you use it?

Answer:

cron is a time-based job scheduler in Linux used to run tasks at specified intervals. It uses the cron daemon to run scheduled commands.

To view/edit the cron jobs for a user, use: `bash crontab -e # Edit the current user's cron jobs` `crontab -l # List current user's cron jobs`

Example of a cron job that runs every day at midnight: `bash 0 0 * * * /path/to/script.sh`

11. What is the purpose of the `/etc/fstab` file?

Answer:

The `/etc/fstab` file contains static information about the system's disk partitions, mount points, and filesystem types. It defines how and where disk partitions, network shares, and other filesystems should be mounted on boot.

Example entry in `/etc/fstab`: `bash /dev/sda1 / ext4 defaults 0 1`

12. What is the difference between `wget` and `curl`?

Answer: - **wget**: A tool used for downloading files over HTTP, HTTPS, or FTP. It is more suited for downloading files recursively or in batch. Example: `bash wget https://example.com/file.tar.gz`

- **curl**: A tool used for transferring data over multiple protocols (HTTP, HTTPS, FTP, etc.) and is often used for interacting with REST APIs. It is more flexible than **wget**. Example: `bash curl -O https://example.com/file.tar.gz`

13. Explain the role of **sysctl** in Linux.

Answer:

sysctl is a tool used to view and modify kernel parameters at runtime. It is commonly used to configure kernel parameters that affect system performance, networking, and security.

Example to view the current value of a kernel parameter: `bash sysctl net.ipv4.ip_forward`

Example to set a new value: `bash sysctl -w net.ipv4.ip_forward=1`

14. How would you find a specific file on a Linux system?

Answer:

The **find** command is used to search for files based on various criteria, such as name, size, or modification time.

Example to find a file by name: `bash find / -name "filename.txt"`

Example to search for a file with a specific extension: `bash find /path/to/search -name "*.log"`

15. What are the different runlevels in Linux?

Answer:

Linux runlevels define the state of the system. The most common runlevels are:
 - **0**: Halt (shutdown) - **1**: Single-user mode (for maintenance) - **3**: Multi-user mode (without GUI) - **5**: Multi-user mode with GUI - **6**: Reboot

To change the runlevel, you use the **init** or **systemctl** command.

16. What is the difference between **service** and **systemctl**?

Answer: - **service**: A legacy command used in older Linux distributions to start, stop, and manage services. - **systemctl**: The command used to control the **systemd** system and service manager, which is now the default in many Linux distributions. **systemctl** provides more control and is used for both service management and system boot processes.

Example with **systemctl**: `bash systemctl start apache2 systemctl stop apache2`

17. How do you monitor disk usage in Linux?

Answer:

To monitor disk usage, you can use commands like `df` and `du`. - `df`: Shows disk space usage for all mounted filesystems. `bash df -h` - `du`: Displays disk usage of a specific directory or file. `bash du -sh /path/to/directory`

18. How can you monitor system logs in Linux?

Answer:

Linux logs system messages, errors, and other information in the `/var/log/` directory. Common log files include: - `/var/log/syslog`: General system messages. - `/var/log/auth.log`: Authentication messages. - `/var/log/dmesg`: Boot and kernel messages.

To view logs, you can use commands like: `bash tail -f /var/log/syslog`
`journalctl` # To view logs managed by `systemd`

These questions cover a variety of key **Linux administration** topics that a **DevOps engineer** with 3 years of experience should be familiar with. They range from file management and process control to network configuration, performance monitoring, and security, which are essential skills for managing Linux systems in a DevOps environment.

Certainly! Here are **additional Linux-based interview questions** that may be relevant for a **DevOps engineer with 3 years of experience**:

19. What is the `grep` command, and how do you use it?

Answer:

The `grep` command is used for searching text using patterns (regular expressions) in files. It prints the lines of a file that match the specified pattern.

Examples: - Search for a word in a file: `bash grep "pattern" file.txt`
- Search recursively in all files in a directory: `bash grep -r "pattern" /path/to/directory` - Ignore case while searching: `bash grep -i "pattern" file.txt`

20. What is `top` command and how is it useful in monitoring system resources?

Answer:

The `top` command provides a dynamic, real-time view of the system's resource usage, including CPU, memory, and process statistics.

- **CPU Usage**: Displays CPU utilization for each process.
- **Memory Usage**: Shows memory usage (physical and swap).

- **Load Average:** Indicates the system load over a period of time (1, 5, 15 minutes).

To run: `bash top`

- Press **q** to quit.
- Press **Shift + P** to sort by CPU usage, **Shift + M** for memory usage.

21. What is the purpose of the `du` command in Linux?

Answer:

The `du` (disk usage) command is used to estimate file space usage. It is commonly used to find out how much disk space a specific directory or file is consuming.

Example to check the size of a directory: `bash du -sh /path/to/directory`

Example to check sizes of all directories: `bash du -h --max-depth=1`

22. What is the significance of `sudo` in Linux?

Answer:

`sudo` (SuperUser Do) allows a permitted user to execute a command as the superuser or another user. It is used to grant temporary administrative privileges to a normal user to execute specific tasks that require root access.

Example: `bash sudo apt-get update` # Runs the update command with root privileges

It uses the configuration file `/etc/sudoers` to manage which users are allowed to execute commands as the superuser.

23. Explain the `/etc/passwd` and `/etc/shadow` files.

Answer: - `/etc/passwd`: Contains information about user accounts, including the username, user ID (UID), group ID (GID), home directory, default shell, and password hash (in earlier Unix versions, the password was stored here, but now it is moved to `/etc/shadow`). - `/etc/shadow`: Stores hashed passwords for users and password-related information (e.g., password expiry, last change, and account expiration). Only root can read this file.

Example line in `/etc/passwd`: `bash user1:x:1001:1001::/home/user1:/bin/bash`

Example line in `/etc/shadow`: `bash user1:6abcd1234$...:17955:0:99999:7:::`

24. How would you create a new user in Linux?

Answer:

To create a new user, you can use the `useradd` command followed by the username.

Example: `bash sudo useradd -m -s /bin/bash newuser`

This creates the user **newuser**, creates their home directory (`-m`), and sets the default shell to `/bin/bash`.

To set a password for the user: `bash sudo passwd newuser`

25. What is a Linux kernel module, and how do you manage it?

Answer:

A **kernel module** is a piece of code that can be loaded or unloaded into the kernel as needed to extend the functionality of the kernel without rebooting the system. Kernel modules are used for device drivers, filesystems, and system calls.

- **Loading a module:** `bash sudo modprobe <module_name>`
- **Unloading a module:** `bash sudo modprobe -r <module_name>`
- **List loaded modules:** `bash lsmod`

26. How do you check the system's uptime and load averages in Linux?

Answer: The system's **uptime** and **load averages** can be checked using the `uptime` or `top` command.

Example using `uptime`: `bash uptime` Output: 15:00:00 up 5 days, 2:35, 3 users, load average: 0.00, 0.01, 0.05

The load averages indicate the system load for the last 1, 5, and 15 minutes.

27. What is the difference between `&&` and `||` in Linux?

Answer:

- **&& (Logical AND):** Executes the second command only if the first command succeeds (returns a status of 0). `bash command1 && command2 # Executes command2 if command1 succeeds` - **|| (Logical OR):** Executes the second command only if the first command fails (returns a non-zero status). `bash command1 || command2 # Executes command2 if command1 fails`

28. What are iptables in Linux?

Answer:

iptables is a user-space utility program that allows a system administrator to configure the **IP packet filter** rules of the Linux kernel firewall. It controls incoming and outgoing network traffic based on defined rules.

Example of listing rules: `bash sudo iptables -L`

Example of blocking incoming traffic on port 80 (HTTP): `bash sudo iptables -A INPUT -p tcp --dport 80 -j DROP`

29. Explain the `ps` and `pstree` commands.

Answer: - **ps:** Displays information about the running processes. Example: `bash ps aux` # Displays detailed information about all running processes

- **pstree:** Displays running processes in a tree-like structure, showing parent-child relationships. Example: `bash pstree` # Displays processes in a tree view

30. What are `systemd` units and how are they used in Linux?

Answer:

systemd is a system and service manager for Linux, and it manages **units**. Units represent resources like services, devices, sockets, mount points, etc., and are configured using unit files located in `/etc/systemd/system/` or `/lib/systemd/system/`.

Types of units: - **Service:** Defines a service process (e.g., a daemon). - **Socket:** Defines communication channels for `systemd` services. - **Target:** Defines groups of `systemd` units that can be started or stopped together.

Example of starting a service: `bash sudo systemctl start apache2`

31. What is the difference between `soft` and `hard` links in Linux?

Answer: - **Hard Link:** A hard link creates a direct reference to the inode of a file, making multiple directory entries that point to the same data. Deleting one hard link doesn't delete the file unless all hard links are deleted. - **Soft (Symbolic) Link:** A symbolic link is a pointer to the file's path, and if the original file is deleted, the symbolic link becomes broken.

Example of creating a symbolic link: `bash ln -s /path/to/original /path/to/symlink`

32. How can you find out what processes are consuming the most memory in Linux?

Answer: You can use the `top` or `htop` command to identify memory-consuming processes: - **top:** Sort processes by memory usage by pressing `Shift + M` in `top`. `bash top` - **htop:** A more user-friendly version of `top` (you may need to install it). `bash sudo apt install htop` `htop`

33. How do you check the free and used memory on a Linux machine?

Answer: - **free** command displays memory usage. `bash free -h #`
Shows memory usage in human-readable format - Alternatively, you can
check memory usage using `/proc/meminfo`: `bash cat /proc/meminfo`

34. What is the `lsof` command in Linux?

Answer:

The `lsof` (List Open Files) command shows which files are open by which processes. It is useful for troubleshooting, especially when identifying processes holding onto files or ports.

Example: `bash lsof -i :80 # Lists processes using port 80`

35. How do you change the hostname of a Linux machine?

Answer:

To change the hostname, you can edit the `/etc/hostname` and `/etc/hosts` files and reboot the machine.

1. Edit the hostname file: `bash sudo nano /etc/hostname`
2. Modify the `/etc/hosts` file to reflect the new hostname: `bash sudo nano /etc/hosts` Change any references to the old hostname to the new one.
3. Reboot the system: `bash sudo reboot`

These additional **Linux-based questions** will test your knowledge on **system administration, networking, and troubleshooting**, which are important areas for a **DevOps engineer**. They also cover practical commands and utilities used for day-to-day operations in a Linux environment.